




Expert-selected evaluation set ($N = 100$)
50 simple | 50 hard questions



Anonymized response pool ($N = 1,300$)
Source labels and identifying cues removed

Response-source-blinded dual-perspective evaluation

Expert panel ($n = 10$)
1,300 ratings; one intended rating per response

Parent panel ($n = 16$)
1,300 ratings; one intended rating per response

Six 1–5 Likert items
Correctness/reliability | Clarity | Completeness
Empathy | Fair and non-misleading | Perceived absence of bias

Primary models: source $\times$ panel + difficulty
Random intercepts: evaluator, question, question-by-source response | Sensitivity: add response length